

Interactive Computer Graphics: Lecture 8

Texture mapping

Some slides adopted from
H. Pfister, Harvard

The Problem:




- We don't want to represent all this detail with geometry

Graphics Lecture 8: Slide 2

The Solution: Textures

- The visual appearance of a graphics scene can be greatly enhanced by the use of texture.
- Consider a brick building, using a polygon for every brick require a huge effort in scene design.
- So why not use one polygon and draw a repeating brick pattern (texture) onto it?

The Quest for Visual Realism



Graphics Lecture 8: Slide 4

Texture Definition

- Textures may be defined as:
 - One-dimensional functions
 - parameter can have arbitrary domain (e.g., incident angle)
 - Two-dimensional functions
 - information is calculated for every (u, v) , many possibilities
 - Raster images (“texels”)
 - Most often used method
 - Images from scanner, photos or calculation
 - Three-dimensional functions
 - Volume $T(u, v, w)$
- Procedural texture vs. raster data

Procedural textures


- Write a function: $F(\mathbf{p}) \rightarrow \text{color}$



- non-intuitive
- difficult to match a texture that already exists in the 'real' world


Graphics Lecture 8: Slide 6

Photo textures




Graphics Lecture 8: Slide 7

The concept of texture mapping




Texture mapping: Terminology




Graphics Lecture 8: Slide 9

The concept of texture mapping



Parametrization

- How to do the mapping?





- Usually objects are not that simple

Graphics Lecture 8: Slide 11

Parametrization: Planar


- Planar mapping: dump one of the coordinates



Only looks good from the front!

Parametrization


- Cylindrical and spherical mapping: compute angles between vertex and object center
- Compare to polar/spherical coordinate systems




Graphics Lecture 8: Slide 13

Parametrization

- Cylindrical and spherical mapping



Cylindrical






Spherical

Graphics Lecture 8: Slide 14

Parametrization: Box

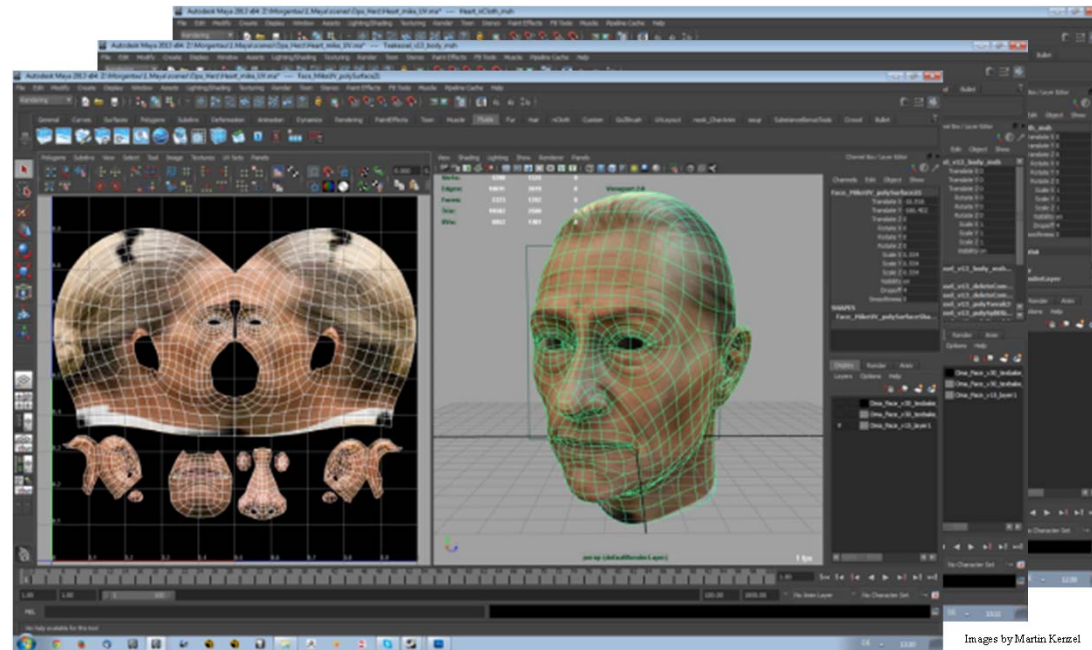
- Box mapping: used mainly for environment mapping (see later)



Graphics Lecture 8: Slide 15

Parametrization

- Manual mapping using CAD Software
 - “Unwrapping”




Images by Martin Kerzel

Graphics Lecture 8: Slide 16

Parameterization problems


- All mappings have distortions and singularities
- Often they need to be fixed manually (CAD software)



Graphics Lecture 8: Slide 17


Texture Coordinates

- Specify a texture coordinate at each vertex
- Canonical texture coordinates $(0,0) \rightarrow (1,1)$
- Often the texture size is a power of 2 (but it doesn't have to be)
- How can we tile this texture?




Texture Addressing

- What happens outside $[0,1]$?
- Border, repeat, clamp, mirror
- Also called texture addressing modes



Texture with red border applied to primitive

Static color




Clamped texture applied to primitive

Clamped





Repeated



Mirrored

Tiling Texture




Texture synthesis



Graphics Lecture 8: Slide 21

Texture Coordinates


- Specify a texture coordinate at each vertex
- Canonical texture coordinates $(0,0) \rightarrow (1,1)$
- Linearly interpolate the values in screen space



Graphics Lecture 8: Slide 22


Mapping texture to individual pixels

- Interpolate texture coordinates across scanlines
- Same as Gouraud shading but now for texture coordinates not shading values



Graphics Lecture 8: Slide :

What goes wrong when we linearly interpolate texture coordinates?




- Notice the distortion along the diagonal triangle edge of the cube face after perspective projection

Graphics Lecture 8: Slide 24

Perspective projection


- The problem is that perspective projection does not preserve linear combinations of points!
- In particular; equal distances in 3D space **do not** map to equal distances in screen space



- Linear interpolation in screen space is not the same as linear interpolation in 3D space


How to fix?

- Assign parameter t to 3-D vertices \mathbf{p} and \mathbf{r}
- t controls linear blend of texture coordinates of \mathbf{p} and \mathbf{r}
- Let $t = 0$ at \mathbf{p} , and $t = 1$ at \mathbf{r}
- Assume for simplicity that the image plane is at $z = 1$ *



How to fix?


- \mathbf{p} projects to \mathbf{p}' and \mathbf{r} projects to \mathbf{r}' (simply divide by z coordinate)
- What value should t have at location \mathbf{q}' ?



Graphics Lecture 8: Slide 27

How to fix?


- We cannot linearly interpolate t between \mathbf{p}' and \mathbf{r}'
- Only projected values can be linearly interpolated in screen space
- Solution: perspective-correct interpolation



Graphics Lecture 8: Slide 28


Perspective Correct Interpolation

- Linearly interpolate t/z (not t) between \mathbf{p}' and \mathbf{r}' .
 - Compute $t_{p'} = t_p / z_p$ $t_{r'} = t_r / z_r$
 - Linearly interpolate (lerp) $t_{p'}$ and $t_{r'}$ to get $t_{q'}$ at location \mathbf{q}'
- But, we want the un-projected parameter t_q (not $t_{q'}$)



Perspective Correct Interpolation


- The parameters t_p & t_q , are related to t_p & t_q by perspective factors of $1/z_p$ and $1/z_q$
 - lerp $1/z_p$ and $1/z_r$ to obtain $1/z_q$ at point q ‘



Graphics Lecture 8: Slide 30

Perspective Correct Interpolation

- The parameters t_p & t_q , are related to t_p & t_q by perspective factors of $1/z_p$ and $1/z_q$
 - lerp $1/z_p$ and $1/z_r$ to obtain $1/z_q$ at point q'
 - Divide $t_{q'}$ by $1/z_q$ to get t_q



Perspective Correct Interpolation

- Summary:


- Given texture parameter t at vertices:

- Compute $1/z$ for each vertex
 - Linearly interpolate $1/z$ across the triangle
 - Linearly interpolate t/z across the triangle
 - Do perspective division:

Divide t/z by $1/z$ to obtain interpolated parameter t


$$t_q = \frac{\text{lerp} \left(\frac{t_p}{z_p}, \frac{t_r}{z_r} \right)}{\text{lerp} \left(\frac{1}{z_p}, \frac{1}{z_r} \right)}$$

What Goes Wrong?




Graphics Lecture 8: Slide 33

Perspective Correct Interpolation



Mapping texture to individual pixels

Alternative



$P_{1..4}$: Polygon vertices
 p : Pixel to be textured


Bilinear Texture mapping

Bi-linear Map - Solving for a and b

$$\mathbf{p} = \alpha \mathbf{a} + \beta \mathbf{e}$$
$$\mathbf{e} = \mathbf{b} + \alpha (\mathbf{c} - \mathbf{b})$$

so

$$\mathbf{p} = \alpha \mathbf{a} + \beta \mathbf{b} + \alpha\beta (\mathbf{c} - \mathbf{b}) \quad \text{Quadratic in the unknowns !}$$



Non Linearities in texture mapping

- The second order term means that straight lines in the texture may become curved when the texture is mapped.
- However, if the mapping is to a parallelogram:

$$\mathbf{p} = \alpha \mathbf{a} + \beta \mathbf{b} + \alpha\beta (\mathbf{c} - \mathbf{b})$$

and

$$\mathbf{b} = \mathbf{c}$$



so

$$\mathbf{p} = \alpha \mathbf{a} + \beta \mathbf{b}$$


Texture Mapping & Illumination

- Texture mapping can be used to alter parts of the illumination equation


$$L(\omega_r) = k_a I_a + (k_d I_d (\mathbf{n} \cdot \mathbf{s}) + k_s I_s (\mathbf{v} \cdot \mathbf{r})^q)$$




Constant Diffuse Color



Texture Image




Texture used as Label



Texture used as Diffuse Color

2D Texture Mapping



- Increases the apparent complexity of simple geometry
- Requires perspective projection correction
- Can specify variations in shading within a primitive:
 - Illumination
 - Surface Reflectance



Graphics Lecture 8: Slide 39

What's Missing?




- What's the difference between a real brick wall and a photograph of the wall texture-mapped onto a plane?
- What happens if we change the lighting or the camera position?




Graphics Lecture 8: Slide 40

Remember Normal Averaging for Shading?


- Instead of using the normal of the triangle, interpolate an averaged normal at each vertex across the face




G




Bump Mapping




- Textures can be used to alter the surface normal of an object.
- Does not change actual shape of the surface – we only shade it as if it were a different shape!



Sphere w/Diffuse Texture



Swirly Bump Map




Sphere w/Diffuse Texture & Bump Map


Graphics Lecture 8: Slide 42

Bump Mapping


- The texture map is treated as a single-valued height function.
- The partial derivatives of the texture tell us how to alter the true surface normal at each point to make the object appear as if it were deformed by the height function.



Cylinder w/Diffuse Texture Map




Bump Map



Cylinder w/Texture Map & Bump Map

Graphics Lecture 8: Slide 43

Another Bump Map Example




Graphics Lecture 8: Slide 44

Image source: Wikipedia, 2016

What's Missing?

- What does a texture- & bump-mapped object look like as you move the viewpoint?
- What does the silhouette of a bump-mapped object look like?




https://threejs.org/examples/webgl_materials_bumpmap.html

Graphics Lecture 8: Slide 45

Image source: Wikipedia, 2016

Displacement Mapping


- Use the texture map to actually move the surface point.
 - How is this different than bump mapping?
- The geometry must be displaced before visibility is determined.



Graphics Lecture 8: Slide 46

Environment Maps

- We can simulate reflections by using the direction of the reflected ray to index a spherical texture map at "infinity".
- Assumes that all reflected rays begin from the same point.



Graphics Lecture 8: Slide 47

Environment Mapping Example



https://threejs.org/examples/webgl_materials_cubemap.html

Graphics Lecture 8: Slide 48

Environment Mapping Example



Graphics Lecture 8: Slide 49

Interactive examples

- https://threejs.org/examples/#webgl_materials_bumpmap
- https://threejs.org/examples/#webgl_materials_displacementmap
- https://threejs.org/examples/webgl_materials_cubemap_dynamic2.html
- <https://www.youtube.com/watch?v=K5n3p97-tuQ>

Interactive Computer Graphics: Lecture 9

Rasterization, Visibility & Anti-aliasing

Some slides adopted from
F. Durand and B. Cutler, MIT

The Graphics Pipeline

Modelling
Transformations

Illumination
(Shading)

Viewing Transformation
(Perspective / Orthographic)


Clipping

Projection
(to Screen Space)

Scan Conversion
(Rasterization)

Visibility / Display

- Rasterizes objects into pixels
- Interpolate values inside objects (color, depth, etc.)



Graphics Lecture 9: Slide 2

The Graphics Pipeline

Modelling
Transformations

Illumination
(Shading)

Viewing Transformation
(Perspective / Orthographic)


Clipping

Projection
(to Screen Space)

Scan Conversion
(Rasterization)

Visibility / Display


- Handles occlusions
- Determines which objects are closest and therefore visible



Graphics Lecture 9: Slide 3

Rasterization


- Determine which pixels are drawn into the framebuffer
- Interpolate parameters (colors, texture coordinates, etc.)



Graphics Lecture 9: Slide 4


Rasterization

- What does interpolation mean?
- Examples: Colors, normals, shading, texture coordinates



Graphics Lecture 9: Slide 5


Coordinate intuition



Graphics Lecture 9: Slide 6

Coordinate intuition


Trilinear coordinates



Graphics Lecture 9: Slide 7

Coordinate intuition


barycentric coordinates



Graphics Lecture 9: Slide 8

Coordinate intuition


barycentric coordinates



Graphics Lecture 9: Slide 9

A triangle in terms of vectors


- We can use vertices **a**, **b** and **c** to specify the three points of a triangle
- We can also compute the edge vectors



Graphics Lecture 9: Slide 10

Points and planes

- The three non-collinear points determine a plane




- Example: The vertices a , b and c determine a plane
- The vectors $b - a$ and $c - a$ form a basis for this plane

Basis vectors

- This (non-orthogonal) basis can be used to specify the location of any point **p** in the plane

$$\mathbf{p} = \mathbf{a} + \beta(\mathbf{b} - \mathbf{a}) + \gamma(\mathbf{c} - \mathbf{a})$$



Barycentric coordinates

- We can reorder the terms of the equation:

$$\begin{aligned}\mathbf{p} &= \mathbf{a} + \beta(\mathbf{b} - \mathbf{a}) + \gamma(\mathbf{c} - \mathbf{a}) \\ &= (1 - \beta - \gamma)\mathbf{a} + \beta\mathbf{b} + \gamma\mathbf{c} \\ &= \alpha\mathbf{a} + \beta\mathbf{b} + \gamma\mathbf{c}\end{aligned}$$

- In other words:

$$\mathbf{p}(\alpha, \beta, \gamma) = \alpha\mathbf{a} + \beta\mathbf{b} + \gamma\mathbf{c}$$

- α , β , γ and called barycentric coordinates

Barycentric coordinates

- **Homogenous barycentric coordinates:**
 - normalised so that $\alpha + \beta + \gamma = \text{area of triangle}$
- **Areal coordinates or absolute barycentric coordinates** : barycentric coordinates *normalized by the area of the original triangle* $\alpha + \beta + \gamma = 1$

Barycentric coordinates

- Barycentric coordinates describe a point **p** as an affine combination of the triangle vertices

$$\mathbf{p}(\alpha, \beta, \gamma) = \alpha \mathbf{a} + \beta \mathbf{b} + \gamma \mathbf{c} \quad \alpha + \beta + \gamma = 1$$

- For any point **p** inside the triangle (**a**, **b**, **c**):

$$0 < \alpha < 1$$


$$0 < \beta < 1$$

$$0 < \gamma < 1$$

- Point on an edge: one coefficient is 0
- Vertex: two coefficients are 0, remaining one is 1

Barycentric coordinates and signed distances


- Let $\mathbf{p} = \alpha\mathbf{a} + \beta\mathbf{b} + \gamma\mathbf{c}$. Each coordinate (e.g. β) is the signed distance from \mathbf{p} to the line through a triangle edge (e.g. \mathbf{ac})



Graphics Lecture 9: Slide 16

Barycentric coordinates and signed distances


- Let $\mathbf{p} = \alpha\mathbf{a} + \beta\mathbf{b} + \gamma\mathbf{c}$. Each coordinate (e.g. β) is the signed distance from \mathbf{p} to the line through a triangle edge (e.g. \mathbf{ac})



Graphics Lecture 9: Slide 17

Barycentric coordinates and signed distances


- Let $\mathbf{p} = \alpha\mathbf{a} + \beta\mathbf{b} + \gamma\mathbf{c}$. Each coordinate (e.g. β) is the signed distance from \mathbf{p} to the line through a triangle edge (e.g. \mathbf{ac})



Graphics Lecture 9: Slide 18

Barycentric coordinates and signed distances

- The signed distance can be computed by evaluating implicit line equations, e.g., $f_{ac}(x,y)$ of edge **ac**



Graphics Lecture 9: Slide 19

Recall: Implicit equation for lines

- Implicit equation in 2D:

$$f(x, y) = 0$$

- Points with $f(x, y) = 0$ are on the line
- Points with $f(x, y) \neq 0$ are not on the line

- General implicit form

$$Ax + By + C = 0$$

- Implicit line through two points (x_a, y_a) and (x_b, y_b)


$$(y_a - y_b)x + (x_b - x_a)y + x_a y_b - x_b y_a = 0$$

Implicit equation for lines: Example

A =

B =

C =




Graphics Lecture 9: Slide 21

Implicit equation for lines: Example

Solution 1: $-2x + 4y = 0$

Solution 2: $2x - 4y = 0$

$$kf(x, y) = 0 \text{ for any } k$$



Graphics Lecture 9: Slide 22


Edge equations

- Given a triangle with vertices (x_a, y_a) , (x_b, y_b) , and (x_c, y_c) .
- The line equations of the edges of the triangle are:

$$f_{ab}(x, y) = (y_a - y_b)x + (x_b - x_a)y + x_a y_b - x_b y_a$$

$$f_{bc}(x, y) = (y_b - y_c)x + (x_c - x_b)y + x_b y_c - x_c y_b$$

$$f_{ca}(x, y) = (y_c - y_a)x + (x_a - x_c)y + x_c y_a - x_a y_c$$



Graphics Lecture 9: Slide 23

Barycentric Coordinates

- Remember that: $f(x,y) = 0 \Leftrightarrow kf(x,y) = 0$
- A barycentric coordinate (e.g. β) is a signed distance from a line (e.g. the line that goes through **ac**)
- For a given point **p**, we would like to compute its barycentric coordinate β using an implicit edge equation.
- We need to choose k such that

$$kf_{ac}(x,y) = \beta$$

Barycentric Coordinates

- We would like to choose k such that: $kf_{ac}(x,y) = \beta$
- We know that $\beta = 1$ at point \mathbf{b} :

$$kf_{ac}(x,y) = 1 \Leftrightarrow k = \frac{1}{f_{ac}(x_b, y_b)}$$

- The barycentric coordinate β for point \mathbf{p} is:

$$\beta = \frac{f_{ac}(x,y)}{f_{ac}(x_b, y_b)}$$

Barycentric Coordinates

- In general, the barycentric area coordinates for point **p** are:


$$\alpha = \frac{f_{bc}(x, y)}{f_{bc}(x_a, y_a)} \quad \beta = \frac{f_{ac}(x, y)}{f_{ac}(x_b, y_b)} \quad \gamma = 1 - \alpha - \beta$$

- Given a point **p** with Cartesian coordinates (x, y) , we can compute its barycentric coordinates (α, β, γ) as above.

Barycentric Coordinates

- In general, the barycentric area coordinates for point **p** are the solution of the linear system of equations:

$$\begin{pmatrix} x_a & x_b & x_c \\ y_a & y_b & y_c \\ 1 & 1 & 1 \end{pmatrix} \begin{pmatrix} \alpha \\ \beta \\ \gamma \end{pmatrix} = \begin{pmatrix} x \\ y \\ 1 \end{pmatrix}$$



Barycentric Coordinates

- Can be easily converted to trilinear coordinates

$P_t(t_1, t_2, t_3)$ in trilinear coordinates has barycentric coordinates of $(t_1 \mathbf{a}, t_2 \mathbf{b}, t_3 \mathbf{c})$

where $\mathbf{a}, \mathbf{b}, \mathbf{c}$, are the side lengths of the triangle.

$P_b(\alpha, \beta, \gamma)$ in barycentric coordinates has trilinear coordinates $(\alpha/\mathbf{a}, \beta/\mathbf{b}, \gamma/\mathbf{c})$

Triangle Rasterization

- Many different ways to generate fragments for a triangle
- Checking (α, β, γ) is one method, e.g.

$$(0 < \alpha < 1 \ \&\& \ 0 < \beta < 1 \ \&\& \ 0 < \gamma < 1)$$

- In practice, the graphics hardware uses optimized methods:
 - fixed point precision (not floating-point)
 - incremental (use results from previous pixel)

Triangle Rasterization


- We can use barycentric coordinates to rasterize and color triangles

```
for all x do
  for all y do
    compute (alpha, beta, gamma) for (x,y)
    if (0 < alpha < 1 and
        0 < beta < 1 and
        0 < gamma < 1 ) then
      c = alpha c0 + beta c1 + gamma c2
      drawpixel(x,y) with color c
```

- The color c varies smoothly within the triangle

Visibility: One triangle

- With one triangle, things are simple
- Pixels never overlap!




Graphics Lecture 9: Slide 31

Hidden Surface Removal

- Idea: keep track of visible surfaces
- Typically, we see only the front-most surface
- Exception: transparency

Visibility: Two triangles


- Things get more complicated with multiple triangles
- Fragments might overlap in screen space!



Graphics Lecture 9: Slide 33

Visibility: Pixels vs Fragments


- Each pixel has a unique framebuffer (image) location
- But multiple fragments may end up at same address




Graphics Lecture 9: Slide 34

Visibility: Which triangle should be drawn first?

- Two possible cases:




green triangle on top




orange triangle on top

Visibility: Which triangle should be drawn first?

- Many other cases possible!



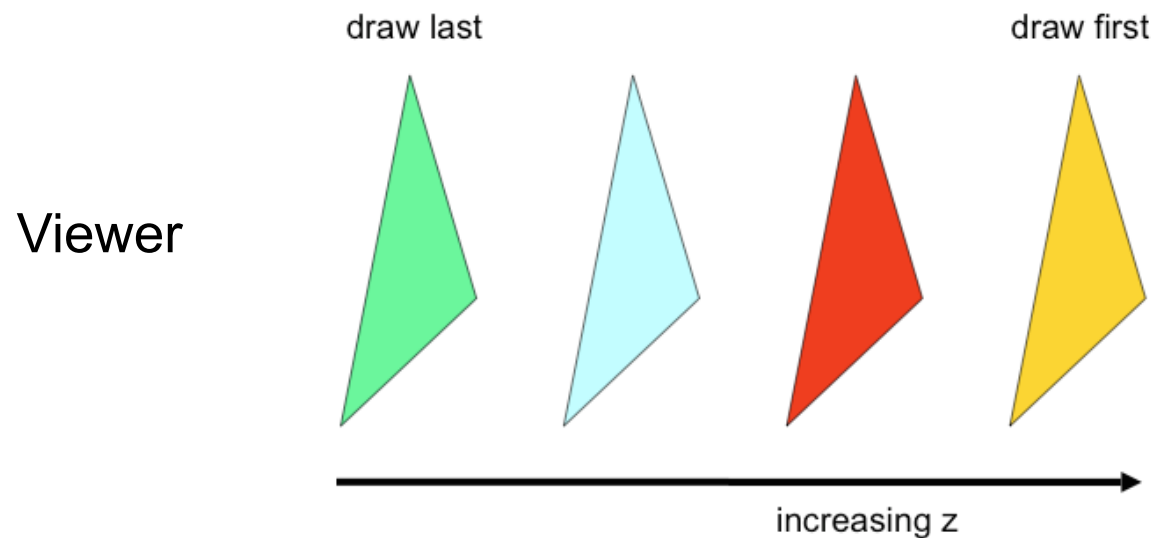
intersection #1



intersection #2

Visibility: Painter's Algorithm


- Sort triangles (using z values in eye space)
- Draw triangles from back to front



Graphics Lecture 9: Slide 37

Visibility: Painter's Algorithm - Problems

- Correctness issues:
 - Intersections
 - Cycles
 - Solve by splitting triangles, but ugly and expensive
- Efficiency (sorting)




Graphics Lecture 9: Slide 38

The Depth Buffer (Z-Buffer)

- Perform hidden surface removal per-fragment
- Idea:
 - Each fragment gets a z value in screen space
 - Keep only the fragment with the smallest z value

The Depth Buffer (Z-Buffer)


- Example:
 - fragment from green triangle has z value of 0.7



Graphics Lecture 9: Slide 40

The Depth Buffer (Z-Buffer)


- Example:
 - fragment from red triangle has z value of 0.3



Graphics Lecture 9: Slide 41

The Depth Buffer (Z-Buffer)

- Since $0.3 < 0.7$, the red fragment wins



Graphics Lecture 9: Slide 42

The Depth Buffer (Z-Buffer)

- Many fragments might map to the same pixel location
- How to track their z-values?
- Solution: z-buffer (2D buffer, same size as image)

1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.1	1.0	1.0	1.0	1.0	1.0
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.1	0.1	1.0	1.0	1.0	1.0	1.0
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.2	0.2	0.3	1.0	1.0	1.0	1.0
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.3	0.3	0.4	1.0	1.0	1.0	1.0
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.3	0.4	0.4	0.5	1.0	1.0	1.0	1.0
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.4	0.4	0.5	0.5	0.5	1.0	1.0	1.0
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.4	0.5	1.0	1.0	1.0

Graphics Lecture 9: Slide 43

The Z-Buffer Algorithm

Let CB be color (frame) buffer, ZB be z-buffer

Initialize z-buffer contents to 1.0 (far away)

For each triangle T

 Rasterize T to generate fragments

 For each fragment F with screen position (x,y,z) and color value C

 If (z < ZB[x,y]) then

 Update color: CB[x,y] = C

 Update depth: ZB[x,y] = z


Z-buffer Algorithm Properties

- What makes this method nice?
 - simple (faciliates hardware implementation)
 - handles intersections
 - handles cycles
 - draw opaque polygons in any order


Alias Effects

- One major problem with rasterization is called alias effects, e.g straight lines or triangle boundaries look jagged
- These are caused by undersampling, and can cause unreal visual artefacts.
- It also occurs in texture mapping

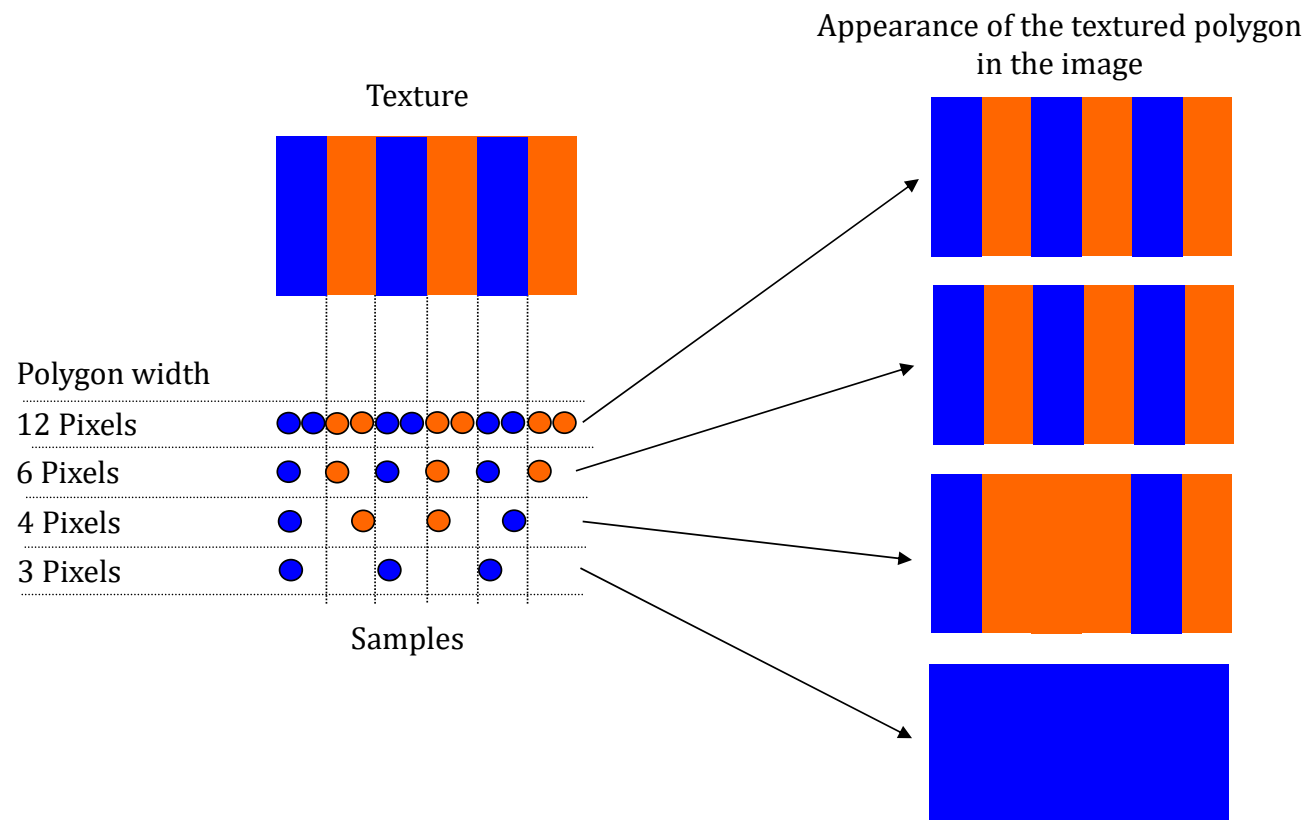
Alias Effects at straight boundaries in raster images.



Desired Boundaries



Pixels Set




Anti-Aliasing

- The solution to aliasing problems is to apply a degree of blurring to the boundary such that the effect is reduced.
- The most successful technique is called **Supersampling**



Supersampling

- The basic idea is to compute the picture at a higher resolution to that of the display area.
- Supersamples are averaged to find the pixel value.
- This has the effect of blurring boundaries, but leaving coherent areas of colour unchanged



Limitations of Supersampling


- Supersampling works well for scenes made up of filled polygons.
- However, it does require a lot of extra computation.
- It does not work for line drawings.




Convolution filtering

- The more common (and much faster) way of dealing with alias effects is to use a 'filter' to blur the image.
- This essentially takes an average over a small region around each pixel

For example consider the image of a line



Treat each pixel of the image





We replace the pixel by a local average,
one possibility would be $3 \cdot I/9$


Weighted averages

- Taking a straight local average has undesirable effects.
- Thus we normally use a weighted average.

$$1/36 *$$

1	4	1
4	16	4
1	4	1





Pros and Cons of Convolution filtering

- Advantages:
 - It is very fast and can be done in hardware
 - Generally applicable
- Disadvantages:
 - It does degrade the image while enhancing its visual appearance.

Anti-Aliasing textures

- Similar
- When we identify a point in the texture map we return an average of texture map around the point.
- Scaling needs to be applied so that the less the samples taken the bigger the local area where averaging is done.

